

```

/*
 * linux/kernel/panic.c
 *
 * Copyright (C) 1991, 1992 Linus Torvalds
 */

/*
 * This function is used through-out the kernel (including mm and fs)
 * to indicate a major problem.
 */
#include <linux/debug_locks.h>
#include <linux/interrupt.h>
#include <linux/kmsg_dump.h>
#include <linux/kallsyms.h>
#include <linux/notifier.h>
#include <linux/module.h>
#include <linux/random.h>
#include <linux/ftrace.h>
#include <linux/reboot.h>
#include <linux/delay.h>
#include <linux/kexec.h>
#include <linux/sched.h>
#include <linux/sysrq.h>
#include <linux/init.h>
#include <linux/nmi.h>

#define PANIC_TIMER_STEP 100
#define PANIC_BLINK_SPD 18

int panic_on_oops = CONFIG_PANIC_ON_OOPS_VALUE;
static unsigned long tainted_mask;
static int pause_on_oops;
static int pause_on_oops_flag;
static DEFINE_SPINLOCK(pause_on_oops_lock);
static bool crash_kexec_post_notifiers;
int panic_on_warn __read_mostly;

int panic_timeout = CONFIG_PANIC_TIMEOUT;
EXPORT_SYMBOL_GPL(panic_timeout);

ATOMIC_NOTIFIER_HEAD(panic_notifier_list);

EXPORT_SYMBOL(panic_notifier_list);

static long no_blink(int state)
{
    return 0;
}

```

```

/* Returns how long it waited in ms */
long (*panic_blink)(int state);
EXPORT_SYMBOL(panic_blink);

/*
 * Stop ourself in panic -- architecture code may override this
 */
void __weak panic_smp_self_stop(void)
{
    while (1)
        cpu_relax();
}

/**
 * panic - halt the system
 * @fmt: The text string to print
 *
 * Display a message, then perform cleanups.
 *
 * This function never returns.
 */
void panic(const char *fmt, ...)
{
    static DEFINE_SPINLOCK(panic_lock);
    static char buf[1024];
    va_list args;
    long i, i_next = 0;
    int state = 0;

    /*
     * Disable local interrupts. This will prevent panic_smp_self_stop
     * from deadlocking the first cpu that invokes the panic, since
     * there is nothing to prevent an interrupt handler (that runs
     * after the panic_lock is acquired) from invoking panic again.
     */
    local_irq_disable();

    /*
     * It's possible to come here directly from a panic-assertion and
     * not have preempt disabled. Some functions called from here want
     * preempt to be disabled. No point enabling it later though...
     */
    /*
     * Only one CPU is allowed to execute the panic code from here. For
     * multiple parallel invocations of panic, all other CPUs either
     * stop themselves or will wait until they are stopped by the 1st CPU
     * with smp_send_stop().
     */
    if (!spin_trylock(&panic_lock))
        panic_smp_self_stop();

```

```

console_verbose();
bust_spinlocks(1);
va_start(args, fmt);
vsnprintf(buf, sizeof(buf), fmt, args);
va_end(args);
pr_emerg("Kernel panic - not syncing: %s\n", buf);
#ifdef CONFIG_DEBUG_BUGVERBOSE
/*
 * Avoid nested stack-dumping if a panic occurs during oops processing
 */
if (!test_taint(TAINT_DIE) && oops_in_progress <= 1)
    dump_stack();
#endif

/*
 * If we have crashed and we have a crash kernel loaded let it handle
 * everything else.
 * If we want to run this after calling panic_notifiers, pass
 * the "crash_kexec_post_notifiers" option to the kernel.
 */
if (!crash_kexec_post_notifiers)
    crash_kexec(NULL);

/*
 * Note smp_send_stop is the usual smp shutdown function, which
 * unfortunately means it may not be hardened to work in a panic
 * situation.
 */
smp_send_stop();

/*
 * Run any panic handlers, including those that might need to
 * add information to the kmsg dump output.
 */
atomic_notifier_call_chain(&panic_notifier_list, 0, buf);

kmsg_dump(KMSG_DUMP_PANIC);

/*
 * If you doubt kdump always works fine in any situation,
 * "crash_kexec_post_notifiers" offers you a chance to run
 * panic_notifiers and dumping kmsg before kdump.
 * Note: since some panic_notifiers can make crashed kernel
 * more unstable, it can increase risks of the kdump failure too.
 */
crash_kexec(NULL);

bust_spinlocks(0);

if (!panic_blink)

```

```

        panic_blink = no_blink;

    if (panic_timeout > 0) {
        /*
         * Delay timeout seconds before rebooting the machine.
         * We can't use the "normal" timers since we just panicked.
         */
        pr_emerg("Rebooting in %d seconds..", panic_timeout);

        for (i = 0; i < panic_timeout * 1000; i += PANIC_TIMER_STEP) {
            touch_nmi_watchdog();
            if (i >= i_next) {
                i += panic_blink(state ^= 1);
                i_next = i + 3600 / PANIC_BLINK_SPD;
            }
            mdelay(PANIC_TIMER_STEP);
        }
    }
    if (panic_timeout != 0) {
        /*
         * This will not be a clean reboot, with everything
         * shutting down. But if there is a chance of
         * rebooting the system it will be rebooted.
         */
        emergency_restart();
    }
#ifdef __sparc__
{
    extern int stop_a_enabled;
    /* Make sure the user can actually press Stop-A (L1-A) */
    stop_a_enabled = 1;
    pr_emerg("Press Stop-A (L1-A) to return to the boot prom\n");
}
#endif
#ifdef CONFIG_S390
{
    unsigned long caller;

    caller = (unsigned long)__builtin_return_address(0);
    disabled_wait(caller);
}
#endif
pr_emerg("---[ end Kernel panic - not syncing: %s\n", buf);
local_irq_enable();
for (i = 0; ; i += PANIC_TIMER_STEP) {
    touch_softlockup_watchdog();
    if (i >= i_next) {
        i += panic_blink(state ^= 1);
        i_next = i + 3600 / PANIC_BLINK_SPD;
    }
}

```

```

        mdelay(PANIC_TIMER_STEP);
    }
}

EXPORT_SYMBOL(panic);

struct tnt {
    u8      bit;
    char    true;
    char    false;
};

static const struct tnt tnts[] = {
    { TAIN_T_PROPRIETARY_MODULE,    'P', 'G' },
    { TAIN_T_FORCED_MODULE,        'F', ' ' },
    { TAIN_T_CPU_OUT_OF_SPEC,      'S', ' ' },
    { TAIN_T_FORCED_RMMOD,         'R', ' ' },
    { TAIN_T_MACHINE_CHECK,        'M', ' ' },
    { TAIN_T_BAD_PAGE,             'B', ' ' },
    { TAIN_T_USER,                 'U', ' ' },
    { TAIN_T_DIE,                  'D', ' ' },
    { TAIN_T_OVERRIDDEN_ACPI_TABLE, 'A', ' ' },
    { TAIN_T_WARN,                 'W', ' ' },
    { TAIN_T_CRAP,                 'C', ' ' },
    { TAIN_T_FIRMWARE_WORKAROUND,  'I', ' ' },
    { TAIN_T_OOT_MODULE,           'O', ' ' },
    { TAIN_T_UNSIGNED_MODULE,      'E', ' ' },
    { TAIN_T_SOFTLOCKUP,           'L', ' ' },
    { TAIN_T_LIVEPATCH,           'K', ' ' },
};

/**
 *      print_tainted - return a string to represent the kernel taint state.
 *
 *      'P' - Proprietary module has been loaded.
 *      'F' - Module has been forcibly loaded.
 *      'S' - SMP with CPUs not designed for SMP.
 *      'R' - User forced a module unload.
 *      'M' - System experienced a machine check exception.
 *      'B' - System has hit bad_page.
 *      'U' - Userspace-defined naughtiness.
 *      'D' - Kernel has oopsed before
 *      'A' - ACPI table overridden.
 *      'W' - Taint on warning.
 *      'C' - modules from drivers/staging are loaded.
 *      'I' - Working around severe firmware bug.
 *      'O' - Out-of-tree module has been loaded.
 *      'E' - Unsigned module has been loaded.
 *      'L' - A soft lockup has previously occurred.

```

```

* 'K' - Kernel has been live patched.
*
* The string is overwritten by the next call to print_tainted().
*/
const char *print_tainted(void)
{
    static char buf[ARRAY_SIZE(tnts) + sizeof("Tainted: ")];

    if (tainted_mask) {
        char *s;
        int i;

        s = buf + sprintf(buf, "Tainted: ");
        for (i = 0; i < ARRAY_SIZE(tnts); i++) {
            const struct tnt *t = &tnts[i];
            *s++ = test_bit(t->bit, &tainted_mask) ?
                    t->>true : t->>false;
        }
        *s = 0;
    } else
        snprintf(buf, sizeof(buf), "Not tainted");

    return buf;
}

int test_taint(unsigned flag)
{
    return test_bit(flag, &tainted_mask);
}
EXPORT_SYMBOL(test_taint);

unsigned long get_taint(void)
{
    return tainted_mask;
}

/**
 * add_taint: add a taint flag if not already set.
 * @flag: one of the TAIN*_ constants.
 * @lockdep_ok: whether lock debugging is still OK.
 *
 * If something bad has gone wrong, you'll want @lockdebug_ok = false, but for
 * some notewortht-but-not-corrupting cases, it can be set to true.
 */
void add_taint(unsigned flag, enum lockdep_ok lockdep_ok)
{
    if (lockdep_ok == LOCKDEP_NOW_UNRELIABLE && __debug_locks_off())
        pr_warn("Disabling lock debugging due to kernel taint\n");

    set_bit(flag, &tainted_mask);
}

```

```

}
EXPORT_SYMBOL(add_taint);

static void spin_msec(int msecs)
{
    int i;

    for (i = 0; i < msecs; i++) {
        touch_nmi_watchdog();
        mdelay(1);
    }
}

/*
 * It just happens that oops_enter() and oops_exit() are identically
 * implemented...
 */
static void do_oops_enter_exit(void)
{
    unsigned long flags;
    static int spin_counter;

    if (!pause_on_oops)
        return;

    spin_lock_irqsave(&pause_on_oops_lock, flags);
    if (pause_on_oops_flag == 0) {
        /* This CPU may now print the oops message */
        pause_on_oops_flag = 1;
    } else {
        /* We need to stall this CPU */
        if (!spin_counter) {
            /* This CPU gets to do the counting */
            spin_counter = pause_on_oops;
            do {
                spin_unlock(&pause_on_oops_lock);
                spin_msec(MSEC_PER_SEC);
                spin_lock(&pause_on_oops_lock);
            } while (--spin_counter);
            pause_on_oops_flag = 0;
        } else {
            /* This CPU waits for a different one */
            while (spin_counter) {
                spin_unlock(&pause_on_oops_lock);
                spin_msec(1);
                spin_lock(&pause_on_oops_lock);
            }
        }
    }
    spin_unlock_irqrestore(&pause_on_oops_lock, flags);
}

```

```

}

/*
 * Return true if the calling CPU is allowed to print oops-related info.
 * This is a bit racy..
 */
int oops_may_print(void)
{
    return pause_on_oops_flag == 0;
}

/*
 * Called when the architecture enters its oops handler, before it prints
 * anything. If this is the first CPU to oops, and it's oopsing the first
 * time then let it proceed.
 *
 * This is all enabled by the pause_on_oops kernel boot option. We do all
 * this to ensure that oopses don't scroll off the screen. It has the
 * side-effect of preventing later-oopsing CPUs from mucking up the display,
 * too.
 *
 * It turns out that the CPU which is allowed to print ends up pausing for
 * the right duration, whereas all the other CPUs pause for twice as long:
 * once in oops_enter(), once in oops_exit().
 */
void oops_enter(void)
{
    tracing_off();
    /* can't trust the integrity of the kernel anymore: */
    debug_locks_off();
    do_oops_enter_exit();
}

/*
 * 64-bit random ID for oopses:
 */
static u64 oops_id;

static int init_oops_id(void)
{
    if (!oops_id)
        get_random_bytes(&oops_id, sizeof(oops_id));
    else
        oops_id++;

    return 0;
}
late_initcall(init_oops_id);

void print_oops_end_marker(void)

```



```

{
    init_oops_id();
    pr_warn("---[ end trace %016llx ]---\n", (unsigned long long)oops_id);
}

/*
 * Called when the architecture exits its oops handler, after printing
 * everything.
 */
void oops_exit(void)
{
    do_oops_enter_exit();
    print_oops_end_marker();
    kmsg_dump(KMSG_DUMP_OOPS);
}

#ifdef WANT_WARN_ON_SLOWPATH
struct slowpath_args {
    const char *fmt;
    va_list args;
};

static void warn_slowpath_common(const char *file, int line, void *caller,
                                unsigned taint, struct slowpath_args *args)
{
    disable_trace_on_warning();

    pr_warn("-----[ cut here ]-----\n");
    pr_warn("WARNING: CPU: %d PID: %d at %s:%d %pS()\n",
            raw_smp_processor_id(), current->pid, file, line, caller);

    if (args)
        vprintk(args->fmt, args->args);

    if (panic_on_warn) {
        /*
         * This thread may hit another WARN() in the panic path.
         * Resetting this prevents additional WARN() from panicking the
         * system on this thread. Other threads are blocked by the
         * panic_mutex in panic().
         */
        panic_on_warn = 0;
        panic("panic_on_warn set ...\n");
    }

    print_modules();
    dump_stack();
    print_oops_end_marker();
    /* Just a warning, don't kill lockdep. */
    add_taint(taint, LOCKDEP_STILL_OK);
}

```

```

}

void warn_slowpath_fmt(const char *file, int line, const char *fmt, ...)
{
    struct slowpath_args args;

    args.fmt = fmt;
    va_start(args.args, fmt);
    warn_slowpath_common(file, line, __builtin_return_address(0),
                        TAIN_T_WARN, &args);
    va_end(args.args);
}
EXPORT_SYMBOL(warn_slowpath_fmt);

void warn_slowpath_fmt_taint(const char *file, int line,
                            unsigned taint, const char *fmt, ...)
{
    struct slowpath_args args;

    args.fmt = fmt;
    va_start(args.args, fmt);
    warn_slowpath_common(file, line, __builtin_return_address(0),
                        taint, &args);
    va_end(args.args);
}
EXPORT_SYMBOL(warn_slowpath_fmt_taint);

void warn_slowpath_null(const char *file, int line)
{
    warn_slowpath_common(file, line, __builtin_return_address(0),
                        TAIN_T_WARN, NULL);
}
EXPORT_SYMBOL(warn_slowpath_null);
#endif

#ifdef CONFIG_CC_STACKPROTECTOR

/*
 * Called when gcc's -fstack-protector feature is used, and
 * gcc detects corruption of the on-stack canary value
 */
__visible void __stack_chk_fail(void)
{
    panic("stack-protector: Kernel stack is corrupted in: %p\n",
        __builtin_return_address(0));
}
EXPORT_SYMBOL(__stack_chk_fail);

#endif

```

```
core_param(panic, panic_timeout, int, 0644);
core_param(pause_on_oops, pause_on_oops, int, 0644);
core_param(panic_on_warn, panic_on_warn, int, 0644);

static int __init setup_crash_kexec_post_notifiers(char *s)
{
    crash_kexec_post_notifiers = true;
    return 0;
}
early_param("crash_kexec_post_notifiers", setup_crash_kexec_post_notifiers);

static int __init oops_setup(char *s)
{
    if (!s)
        return -EINVAL;
    if (!strcmp(s, "panic"))
        panic_on_oops = 1;
    return 0;
}
early_param("oops", oops_setup);
```